

Data Science & ML

Case Study 2

Unsupported Cell Type. Double-Click to inspect/edit the content.

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
# Load the dataset cereal.csv.
#data = pd.read_csv('cereal.csv')
data = pd.read_csv('/content/drive/My Drive/Colab_Data/cereal.csv')
```

```
print("Info of the dataset:")
print(data.info())
print("-"*50)
print("Check for null values in the dataset:")
print(data.isnull().sum())
print("-"*50)
print("Print the first rows from the dataset:")
print(data.head())
```

```
<Class 'pandas.core.frame.DataFrame'>
RangeIndex: 77 entries, 0 to 76
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   name        77 non-null    object
1   mfr         77 non-null    object
2   type        77 non-null    object
3   calories    77 non-null    int64
4   protein     77 non-null    int64
5   fat         77 non-null    int64
6   sodium      77 non-null    int64
7   fiber       77 non-null    float64
8   carbo       77 non-null    float64
9   sugars      77 non-null    int64
10  potass      77 non-null    int64
11  vitamins    77 non-null    int64
12  shelf       77 non-null    int64
13  weight      77 non-null    float64
14  cups        77 non-null    float64
15  rating      77 non-null    float64
dtypes: float64(5), int64(8), object(3)
memory usage: 9.8+ KB
None
```

Check for null values in the dataset:

```
name      0
mfr       0
type      0
calories  0
protein   0
fat       0
sodium    0
fiber     0
carbo     0
sugars    0
potass    0
vitamins  0
shelf     0
weight    0
cups      0
rating    0
```

```

name mfr type calories protein fat sodium fiber \
0 100% Bran N C 70 4 1 130 10.0
1 100% Natural Bran Q C 120 3 5 15 2.0
2 All-Bran K C 70 4 1 260 9.0
3 All-Bran with Extra Fiber K C 50 4 0 140 14.0
4 Almond Delight R C 110 2 2 200 1.0

carbo sugars potass vitamins shelf weight cups rating
0 5.0 6 280 25 3 1.0 0.33 68.402973
1 8.0 8 135 0 3 1.0 1.00 33.983679
2 7.0 5 320 25 3 1.0 0.33 59.425505
3 8.0 0 330 25 3 1.0 0.50 93.704912
4 14.0 8 -1 25 3 1.0 0.75 34.384843

```

Unsupported Cell Type. Double-Click to inspect/edit the content.

```

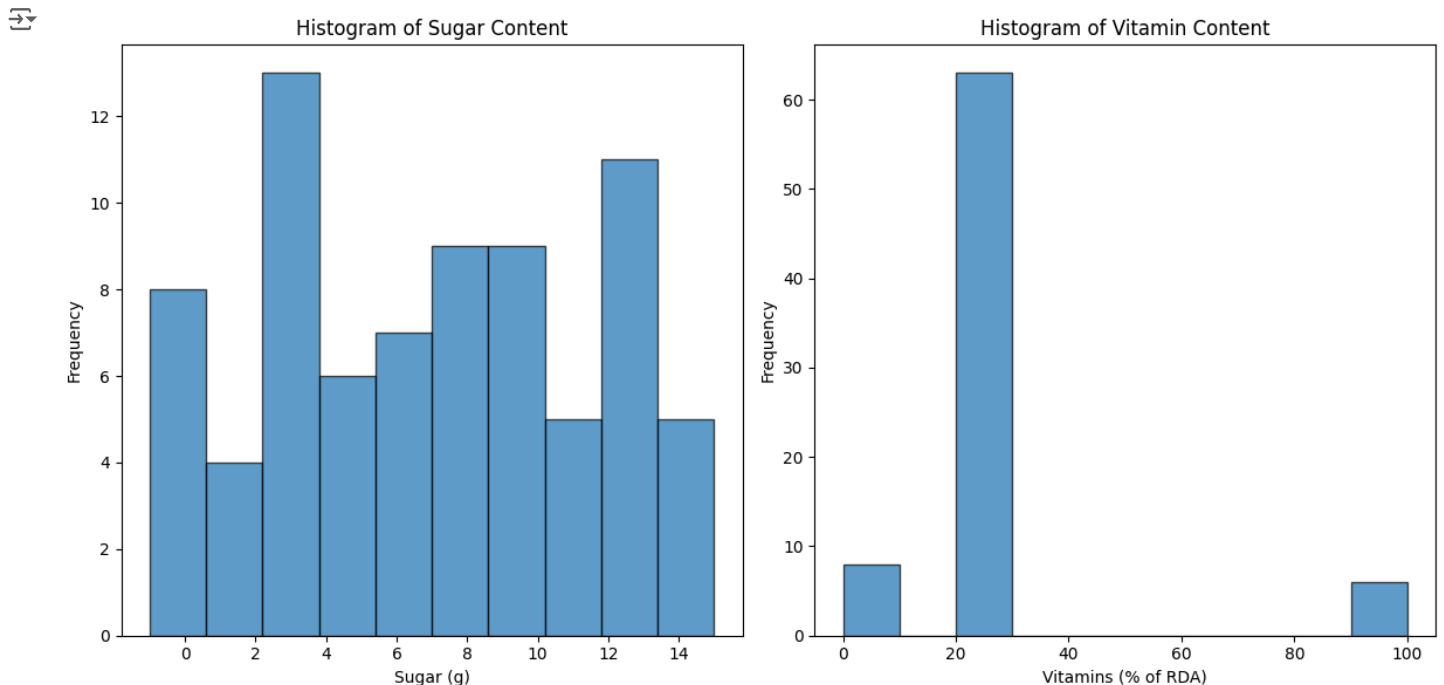
#1
# Plot histograms
plt.figure(figsize=(12, 6))

# Histogram for sugar content
plt.subplot(1, 2, 1)
plt.hist(data['sugars'], bins=10, edgecolor='black', alpha=0.7)
plt.title('Histogram of Sugar Content')
plt.xlabel('Sugar (g)')
plt.ylabel('Frequency')

# Histogram for vitamin content
plt.subplot(1, 2, 2)
plt.hist(data['vitamins'], bins=10, edgecolor='black', alpha=0.7)
plt.title('Histogram of Vitamin Content')
plt.xlabel('Vitamins (% of RDA)')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()

```



2. Create a new column with full manufacturer names and plot a bar chart

```

# Mapping manufacturer codes to full names
manufacturer_mapping = {
    'N': 'Nabisco',
    'Q': 'Quaker Oats',
    'K': 'Kelloggs',
    'R': 'Raslston Purina',
    'G': 'General Mills',

```

```

    'P': 'Post',
    'A': 'American Home Foods Products'
}

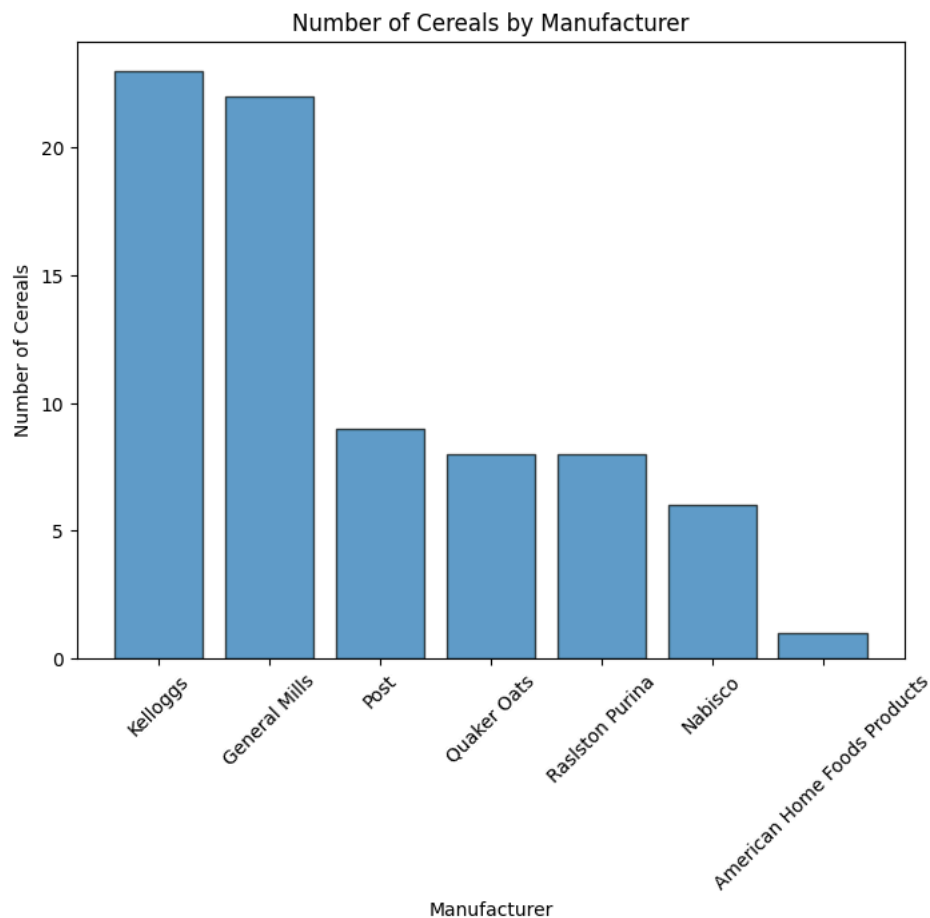
# Add a new column with full manufacturer names
data['manufacturer_full'] = data['mfr'].map(manufacturer_mapping)

# Count the number of cereals by manufacturer
manufacturer_counts = data['manufacturer_full'].value_counts()

# Plot the bar chart
plt.figure(figsize=(8, 6))
plt.bar(manufacturer_counts.index, manufacturer_counts.values, alpha=0.7, edgecolor='black')
plt.title('Number of Cereals by Manufacturer')
plt.xlabel('Manufacturer')
plt.ylabel('Number of Cereals')
plt.xticks(rotation=45)
plt.show()

print(manufacturer_counts)

```



manufacturer_full

Kelloggs	23
General Mills	22
Post	9
Quaker Oats	8
Ralston Purina	8
Nabisco	6
American Home Foods Products	1

```
# 3. Extract target and predictors, split the data
```

```

# Import necessary libraries
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import seaborn as sns
import matplotlib.pyplot as plt

```

```

# Step 1. Define the target variable (y) and predictors (x)
# The 'rating' column is the target, and the predictors are all numeric columns except 'rating'.
y = data['rating']
x = data.select_dtypes(include=['int64', 'float64']).drop(columns=['rating'])

# Step 2. Compute the correlation matrix for feature selection
correlation_matrix = x.corr()

# Function to drop highly correlated features
def drop_highly_correlated(corr_matrix, threshold=0.9):
    """
    Identifies and returns features to drop based on a correlation threshold.
    Only the lower triangle of the correlation matrix is scanned to avoid redundancy.
    """
    to_drop = set() # Set to keep track of features to drop
    cols = corr_matrix.columns

    # Iterate over the lower triangle of the matrix
    for i in range(len(cols)):
        for j in range(i): # Only check lower triangle (j < i)
            if abs(corr_matrix.iloc[i, j]) > threshold: # Check if correlation exceeds the threshold
                print(f"Highly correlated pair: {cols[i]} and {cols[j]} with correlation {abs(corr_matrix.iloc[i, j])}")
                to_drop.add(cols[i]) # Add one of the features to the set

    return list(to_drop)

# Step 3. Visualize the correlation matrix as a heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)
plt.title("Correlation Heatmap of Independent Variables")
plt.show()

# Identify redundant features with a threshold of 0.68
redundant_features = drop_highly_correlated(correlation_matrix, threshold=0.68)
print(f"Features to drop: {redundant_features}")


# Drop redundant features from the dataset
x_reduced = x.drop(columns=redundant_features)

#Step 4. Split the data into training and testing sets
# 75% of the data is used for training, and 25% for testing
x_train, x_test, y_train, y_test = train_test_split(x_reduced, y, test_size=0.25, random_state=42)

print("Training data shape:", x_train.shape)
print("Test data shape:", x_test.shape)

# Display a preview of the training data
print(x_train.head())

```



Highly correlated pair: potass and fiber with correlation 0.9033736685942043

Highly correlated pair: weight and calories with correlation 0.6960910769169041

Features to drop: ['potass', 'weight']

Training data shape: (57, 10)

Test data shape: (20, 10)

	calories	protein	fat	sodium	fiber	carbo	sugars	vitamins	shelf	\
30	100	2	0	45	0.0	11.0	15	25	1	
40	110	2	1	260	0.0	21.0	3	25	2	
39	140	3	1	170	2.0	20.0	9	100	3	
16	100	2	0	290	1.0	21.0	2	25	1	
65	90	3	0	0	3.0	20.0	0	0	1	

cups

30 0.88

40 1.50

39 0.75

16 1.00

65 0.67

4. Fit a linear regression model and evaluate mean squared error

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import cross_val_score
```

Step 1: Fit a linear regression model and evaluate performance

Initialize the linear regression model

```
model = LinearRegression()
```

Train the model on the training data

```
model.fit(x_train, y_train)
```

Display the coefficients of the linear regression model

```
print("Coefficients: ", model.coef_)
```

Make predictions on the test data


```
y_pred = model.predict(x_test)
```

Calculate the Mean Squared Error (MSE) on the test data

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f"Mean Squared Error on Test Data: {mse:.2f}")
```

```
# Step 2: Perform Cross-Validation
# Evaluate the model using 5-fold cross-validation
cv_scores = cross_val_score(model, x_reduced, y, cv=5, scoring='r2')
print(f"Cross-Validation R² Scores: {cv_scores}")
print(f"Mean CV R²: {cv_scores.mean():.4f}")
```

 Coefficients: [-0.22140948 2.98332938 -2.07671607 -0.05408762 2.61721877 1.05226752
 -0.80333669 -0.05374602 -0.13755986 0.63127598]
 Mean Squared Error on Test Data: 1.15
 Cross-Validation R² Scores: [0.99545177 0.99201605 0.9913534 0.98657617 0.99529792]
 Mean CV R²: 0.9921

✓ 1. Mean Squared Error (MSE)

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Where:

- n : Number of data points.
- y_i : Actual value of the target variable for the i th observation.
- \hat{y}_i : Predicted value by the model for the i th observation.

Meaning:

- MSE is the average squared difference between the actual and predicted values.
- It penalizes larger errors more heavily due to squaring.
- A lower MSE indicates better model performance, with 0 being the ideal (perfect predictions).

Interpretation:

- MSE is in the same units as the square of the target variable.
- For example, if rating is measured in points, MSE is in points squared.
- While useful for comparison, it lacks intuitive interpretability since it doesn't represent actual error magnitude.

2. Coefficient of Determination (R^2)

$$R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$$

Where:

- $SS_{\text{res}} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$: Residual sum of squares (unexplained variance).
- $SS_{\text{tot}} = \sum_{i=1}^n (y_i - \bar{y})^2$: Total sum of squares (variance in the target variable).
- \bar{y} : Mean of the actual target values.

Meaning:

- R^2 measures the proportion of variance in the target variable explained by the model.
- Values range from 0 to 1:
- $R^2 = 1$: Perfect fit; the model explains all variance.
- $R^2 = 0$: The model explains no variance (equivalent to predicting the mean of y).
- Negative R^2 : The model performs worse than simply predicting the mean.

Interpretation:

- $R^2 = 0.99$: The model explains 99% of the variance in the target variable.
- While R^2 is useful, it does not penalize overfitting and may appear high for overly complex models.