

# Use-Case-3

December 1, 2024

## 1 Module 8: Dimensionality Reduction

### 1.1 Case Study – 3

#### 1.1.1 Dimensionality Reduction and Supervised Learning for Breast Cancer Classification: A Comparative Study of PCA, LDA, and Ensemble Methods.

```
[8]: # Step 0: Import required libraries
# Load and Explore the Data
import pandas as pd
import numpy as np
# 1. Load the digits dataset breast-cancer-data.csv
dataset=pd.read_csv('breast-cancer-data.csv')
dataset.info()
print(dataset.isnull().sum())
dataset.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	object
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64
16	smoothness_se	569 non-null	float64

17	compactness_se	569	non-null	float64
18	concavity_se	569	non-null	float64
19	concave points_se	569	non-null	float64
20	symmetry_se	569	non-null	float64
21	fractal_dimension_se	569	non-null	float64
22	radius_worst	569	non-null	float64
23	texture_worst	569	non-null	float64
24	perimeter_worst	569	non-null	float64
25	area_worst	569	non-null	float64
26	smoothness_worst	569	non-null	float64
27	compactness_worst	569	non-null	float64
28	concavity_worst	569	non-null	float64
29	concave points_worst	569	non-null	float64
30	symmetry_worst	569	non-null	float64
31	fractal_dimension_worst	569	non-null	float64

dtypes: float64(30), int64(1), object(1)

memory usage: 142.4+ KB

id	0
diagnosis	0
radius_mean	0
texture_mean	0
perimeter_mean	0
area_mean	0
smoothness_mean	0
compactness_mean	0
concavity_mean	0
concave points_mean	0
symmetry_mean	0
fractal_dimension_mean	0
radius_se	0
texture_se	0
perimeter_se	0
area_se	0
smoothness_se	0
compactness_se	0
concavity_se	0
concave points_se	0
symmetry_se	0
fractal_dimension_se	0
radius_worst	0
texture_worst	0
perimeter_worst	0
area_worst	0
smoothness_worst	0
compactness_worst	0
concavity_worst	0
concave points_worst	0
symmetry_worst	0

```
fractal_dimension_worst    0
dtype: int64
```

```
[8]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

  

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

  

	...	radius_worst	texture_worst	perimeter_worst	area_worst	\
0	...	25.38	17.33	184.60	2019.0	
1	...	24.99	23.41	158.80	1956.0	
2	...	23.57	25.53	152.50	1709.0	
3	...	14.91	26.50	98.87	567.7	
4	...	22.54	16.67	152.20	1575.0	

  

	smoothness_worst	compactness_worst	concavity_worst	concave points_worst	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	
2	0.1444	0.4245	0.4504	0.2430	
3	0.2098	0.8663	0.6869	0.2575	
4	0.1374	0.2050	0.4000	0.1625	

  

	symmetry_worst	fractal_dimension_worst
0	0.4601	0.11890
1	0.2750	0.08902
2	0.3613	0.08758
3	0.6638	0.17300
4	0.2364	0.07678

[5 rows x 32 columns]

```
[2]: # Step 1:
# Dropping irrelevant columns like id.
# Encoding categorical labels (diagnosis) into numeric form.
# Standardizing the data for dimensionality reduction.

# Drop 'id' column
dataset = dataset.drop(columns=['id'])
```

```

# Encode 'diagnosis' column (M -> 1, B -> 0)
dataset['diagnosis'] = dataset['diagnosis'].map({'M': 1, 'B': 0})

# Separate features and target
X = dataset.drop(columns=['diagnosis'])
y = dataset['diagnosis']

# Standardize the features
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_standardized = scaler.fit_transform(X)

```

```

[3]: # Step 2: Apply Dimensionality Reduction (PCA & LDA)
# PCA:
# Use PCA to reduce dimensions while retaining a high percentage of variance (e.
#   ↳ g., 95%).

from sklearn.decomposition import PCA

# Perform PCA
pca = PCA(n_components=0.95)
X_pca = pca.fit_transform(X_standardized)

# Check the number of components retained
num_components_pca = pca.n_components_
print(f"Number of components retained by PCA: {num_components_pca}")

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

# Perform LDA
lda = LDA(n_components=1) # Since it's a binary classification problem
X_lda = lda.fit_transform(X_standardized, y)
print(f"LDA reduced the data to {X_lda.shape[1]} component(s).")

# LDA is a supervised dimensionality reduction technique, and the maximum
#   ↳ number of components (k) is determined by the number of classes in the
#   ↳ dataset.
# The formula is:
#   max = C - 1
# Where C is the number of unique classes.
# Since we are working with a binary classification problem (classes: Malignant
#   ↳ and Benign, C=2), the maximum k is 1.

```

```
# This is why we reduced the data to 1 component in LDA.
# LDA ensures that this single component contains the most discriminative
  ↳ information for separating the two classes.
```

Number of components retained by PCA: 10

LDA reduced the data to 1 component(s).

```
[4]: # Step 3: Compare Models with and Without Dimensionality Reduction
# Train and evaluate logistic regression models:

# Without dimensionality reduction.
# After PCA.
# After LDA.

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix,
  ↳ classification_report

# Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X_standardized, y,
  ↳ test_size=0.2, random_state=42)
X_train_pca, X_test_pca = train_test_split(X_pca, test_size=0.2,
  ↳ random_state=42)
X_train_lda, X_test_lda = train_test_split(X_lda, test_size=0.2,
  ↳ random_state=42)

# Initialize Logistic Regression
logistic_model = LogisticRegression(max_iter=10000)

# Baseline Model (No Dimensionality Reduction)
logistic_model.fit(X_train, y_train)
y_pred_baseline = logistic_model.predict(X_test)
accuracy_baseline = accuracy_score(y_test, y_pred_baseline)
print(f"Baseline Model Accuracy: {accuracy_baseline * 100:.2f}%")

# PCA Model
logistic_model.fit(X_train_pca, y_train)
y_pred_pca = logistic_model.predict(X_test_pca)
accuracy_pca = accuracy_score(y_test, y_pred_pca)
print(f"PCA Model Accuracy: {accuracy_pca * 100:.2f}%")

# LDA Model
logistic_model.fit(X_train_lda, y_train)
y_pred_lda = logistic_model.predict(X_test_lda)
accuracy_lda = accuracy_score(y_test, y_pred_lda)
print(f"LDA Model Accuracy: {accuracy_lda * 100:.2f}%")
```

Baseline Model Accuracy: 97.37%  
PCA Model Accuracy: 98.25%  
LDA Model Accuracy: 97.37%

```
[5]: # Step 4: Evaluate Model Performance
# Evaluate performance for each model using confusion matrix, precision, recall, and F1-score.

# Function to evaluate model
def evaluate_model(y_true, y_pred, model_name):
    print(f"\n{model_name} Model Evaluation:")
    print("Confusion Matrix:")
    print(confusion_matrix(y_true, y_pred))
    print("\nClassification Report:")
    print(classification_report(y_true, y_pred))

# Evaluate each model
evaluate_model(y_test, y_pred_baseline, "Baseline")
evaluate_model(y_test, y_pred_pca, "PCA")
evaluate_model(y_test, y_pred_lda, "LDA")
```

Baseline Model Evaluation:

Confusion Matrix:

```
[[70  1]
 [ 2 41]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	71
1	0.98	0.95	0.96	43
accuracy			0.97	114
macro avg	0.97	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114

PCA Model Evaluation:

Confusion Matrix:

```
[[70  1]
 [ 1 42]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	71
1	0.98	0.98	0.98	43

accuracy			0.98	114
macro avg	0.98	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

LDA Model Evaluation:

Confusion Matrix:

```
[[70  1]
 [ 2 41]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	71
1	0.98	0.95	0.96	43
accuracy			0.97	114
macro avg	0.97	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114

[ ]:

Summary and Conclusion Based on the results:

1. Baseline Model Accuracy: 97.37Confusion Matrix: True Negatives: 70 False Positives: 1 False Negatives: 2 True Positives: 41 Precision, Recall, and F1-Score: Class 0 (No Cancer): Precision: 97Class 1 (Cancer): Precision: 98Observations: High accuracy but uses all 30 features, making the model complex and less interpretable.

2. PCA Model Number of Components Retained: 10 Accuracy: 98.25Confusion Matrix: True Negatives: 70 False Positives: 1 False Negatives: 1 True Positives: 42 Precision, Recall, and F1-Score: Class 0 (No Cancer): Precision: 99Class 1 (Cancer): Precision: 98Observations: PCA achieved the highest accuracy with only 10 components, significantly reducing dimensionality while improving accuracy. Superior precision and recall for both classes compared to the baseline.

3. LDA Model Number of Components Retained: 1 Accuracy: 97.37Confusion Matrix: True Negatives: 70 False Positives: 1 False Negatives: 2 True Positives: 41 Precision, Recall, and F1-Score: Class 0 (No Cancer): Precision: 97Class 1 (Cancer): Precision: 98Observations: LDA simplifies the model to a single component, achieving the same accuracy as the baseline. Excellent interpretability and reduced complexity make it ideal for communicating results to doctors.

Final Recommendation PCA: The best performing model with an accuracy of 98.25LDA: While its accuracy matches the baseline, the simplicity of reducing to just 1 component and its focus on class separability make it a highly interpretable choice for classification tasks. Baseline: High accuracy, but the complexity of using all 30 features makes it less practical for deployment or explanation.

Conclusion: Use PCA for the best accuracy with reduced dimensions. Consider LDA for situations where simplicity and interpretability are the primary goals.PCA and LDA: Designed primarily for

dimensionality reduction. PCA is versatile (works for both regression and classification). LDA is tailored for classification, especially when interpretability matters.

DT and RF: Naturally handle feature selection, so dimensionality reduction is often not required. Work well on raw datasets, even with irrelevant features.

Practical Usage: Use PCA or LDA for reducing dimensions before applying regression models or simpler classification algorithms like logistic regression. Use DT or RF directly on raw data unless dimensionality is extremely high, in which case PCA or LDA can still aid preprocessing.

[ ]:

[ ]: